

KEVIN DOSHI

Austin, TX | doshikevin95@gmail.com | (732) 983 1042 | [LinkedIn](#) | [GitHub](#)

PROFESSIONAL EXPERIENCE

Hewlett Packard Enterprise Company

Principal Data Scientist | Austin, Texas

Apr 2026 – Present

- Leading 4-person cross-functional technical team building a trade compliance analytics platform processing 1B+ monthly transactions, with operational and executive dashboards to reduce compliance review time across global operations
- Architecting an end-to-end revenue forecasting product, owning time-series model design, data pipelines, baseline forecast generation, and validation to support long-term financial planning across the product portfolio
- Building scenario and sensitivity modeling capabilities on top of the forecasting system, translating model outputs into actionable storytelling for executive pricing decisions and portfolio-level strategic planning
- Leading causal impact validation of pricing policy change using Difference-in-Differences design, guiding junior data scientist through methodology and analysis to isolate treatment effect, confirming 5% improvement in win rate
- Built a GenAI-powered topic modeling and sentiment pipeline (RAG) on support tickets at scale, delivering structured insights to Finance and leadership that reduced issue resolution time

Senior Data Scientist | Austin, Texas

Jan 2022 – Mar 2026

- Built and deployed a production agentic AI solution with multi-agent orchestration (OpenAI LLMs, LangChain, RAG) enabling non-technical users to autonomously query data through natural language, reducing team request volume by 40%
- Engineered a Responsible AI validation agent within the multi-agent pipeline to detect hallucinations and query vulnerabilities in LLM-generated outputs, establishing safety guardrails and ensuring reliable, enterprise-grade responses at scale
- Built and deployed a propensity-to-buy model on 100M+ records (Python, HiveQL) in partnership with Finance and Product teams, engineering a scalable feature pipeline to score distributor likelihood of purchase improved pricing accuracy by 7%
- Designed and executed an A/B experiment to validate a propensity-to-buy model built and measuring causal impact on revenue outcomes, resulting in 3% revenue uplift across the product portfolio
- Led GTM performance analysis of distributor and reseller purchasing patterns across global geographies, identifying process gaps and translating channel behavior into recommendations that improved partner program operations, deal quality, and pricing accuracy across the product portfolio
- Built and causally validated (Regression Discontinuity design) an escalation threshold optimization model using clustering and margin-constrained optimization, reducing unnecessary escalations and improving Finance review efficiency by 15%
- Built and deployed time series forecasting models (Prophet, ARIMA) to predict next-quarter sell-out for SMB customers, enabling sales team to prioritize outreach proactively; validated via backtesting with 80% forecast accuracy
- Developed churn prediction model (Snowflake, Snowpark) on 500K+ customer records achieving 85% precision, enabling sales and product stakeholders to prioritize proactive retention outreach and protect revenue
- Owned end-to-end design and delivery of a self-serve forecasting platform (Python, SQL, Django front-end) for Finance leadership, automating Budget-vs-Actual variance analysis and monthly forecasting workflows cutting manual prep time by 67%
- Mentored 8 junior data scientists and led capability-building initiatives to upskill analysts on ML best practices, analytics methodologies, and governed workflows coordinating project assignments to maximize team impact and professional growth

Data Scientist III | Roseville, California

Jan 2020 – Dec 2021

- Developed ML risk assessment model in Python analyzing transactional patterns across 5,000 monthly orders, enabling focused review of 250 highest-risk cases and improving operational efficiency for review teams

IT Engineer – Analytics | Santa Clara, California

Sep 2018 – Jan 2020

- Designed and automated dashboards and data pipelines (Python, SQL, Power BI) across IT initiatives providing senior leadership with visibility and enabling data-driven operational decisions improving operational efficiency by 16%

Data Engineering & Analytics Intern | Santa Clara, California

May 2018 – Aug 2018

- Designed database architecture for RF test data; remodeled a legacy analysis tool (SQL, Python, VBA) reducing processing time 30%; and implemented regression algorithms to reduce test flows and automate testing, cutting testing time 20%.

TECHNICAL SKILLS

Analytical & Modeling Techniques: Causal Inference, Propensity Modeling, Customer Segmentation, Time Series Forecasting, Regression, Classification, Clustering, A/B Testing, Hypothesis Testing

Data Engineering & Pipelines: Python (ETL), REST APIs, GraphQL, Snowflake, Snowpark, Vertica, HiveQL, Presto, SQL Server

GenAI & Agentic Frameworks: LangChain, OpenAI LLMs, RAG, FAISS, Prompt Engineering, Fine-tuning

Visualization & Reporting: Tableau, Power BI, Qlik Sense, Dataiku

EDUCATION & CERTIFICATIONS

Applied Generative AI Specialization - Purdue University

Jun 2025

M.S. Information Technology & Analytics - Rutgers University, NJ, US

GPA: 3.80/4.0 | Jan 2019

B.E. Information Technology - University of Mumbai, India

GPA: 3.73/4.0 | May 2017